

Reinforcement Learning Basics


Weizi Li

Department of Computer Science
University of Memphis



- What is RL
- Reward
- RL Agent
- Brief History of RL

What is RL



- 2015: Atari games **superhuman on some games**
- 2016: Game of Go (AlphaGo) **superhuman**
- 2017: Game of Go and Chess (AlphaZero) **superhuman**
- 2018: Robot manipulation (real robots)
- 2018: Mixed-autonomy traffic
- 2019: Dota 2 (OpenAI Five) **superhuman**
- 2019: StarCraft II (AlphaStar) **grandmaster**
- 2019: Solving Rubik's cube with a robot hand (real robots)
- 2021: Atari games (Go-Explore) **superhuman on all games**
- 2021: Chip floor planning **superhuman**

- How an autonomous agent that senses and acts in its environment can learn to choose optimal actions sequentially to achieve its goal

- Computer science: games, robotics...
- Operations: inventories, production, revenue management...
 - ▶ Manufacturing: choose a sequence of actions to maximize profits.
 - ▶ Taxi scheduling: pick up passengers in a city to minimize waiting time.
- Finance: portfolio management, options...
- Control: process control, path planning, robotics...

- The world that the agent lives in
 - ▶ fully observable vs. partially observable
 - ▶ single-agent vs. multiagent
 - ▶ deterministic vs. nondeterministic
 - ▶ known vs. unknown

- Agent can sense the environment and execute actions in the environment

- The movement or other types of decisions available to the agent
- An executed action will affect the environment
- An action may have a long-term consequence
- Choosing an optimal action is difficult: make decisions beyond the immediate payoff

- A snapshot of the environment at a particular point in time
- Fully observed: agent observes environment fully
- Partially observed: agent observes environment partially (e.g., robot with a front-facing camera)
- May contain irrelevant information for a specific task
- Information in state may be perturbed due to flaws of sensor/processing algorithms


- A numerical value (scalar)
- An immediate payoff to each distinct action that agent chooses to take in each distinct state
- Can be 0 (no reward) or even -1 (penalty) at intermediate steps of many tasks
- The actual reward can be *sparse* and *delayed*

- A function executed by the environment to determine the immediate payoff after agent executes an action in a state


- A function executed by the environment to determine the next state after agent executes an action in the current state

- Defines the behavior of an agent by mapping states to actions
- The optimal policy is one that, from any initial state, chooses actions that maximize the reward accumulated over time

- Agent (via a policy) chooses an action to execute in the current state



- Environment launches its reward function and transition function to produce immediate payoff to agent and transit agent to the next state



- Perform sequences of actions, observe their consequences (i.e., rewards), and learn an optimal policy
- During learning, agent does not need any prior knowledge of the effects of its actions on the environment

- Supervised learning: recognition problem
- RL: decision problem
- Learn via trial and error, no supervisor only indirect, delayed feedback
- Action order (time) matters, non-i.i.d. data
- More difficult learning task: usually learn fewer parameters (dozes and thousands) than supervised learning (millions) due to limited feedback

- Agent
- Action
- Environment
- State
- Reward



- Agent
- Action
- Environment
- State
- Reward



- Agent
- Action
- Environment
- State
- Reward



- Agent
- Action
- Environment
- State
- Reward



Reward

- “reward / immediate reward / indirect reward”: immediate payoff at each timestep, could be “no reward” or “penalty”

- “cumulative reward”: sum of reward at each timestep

- “delayed reward / sparse reward”: reward you actually want

- The goal of RL agent is maximizing expected cumulative reward

- Manually designed reward can be brittle:
[https://deepmind.com/blog/article/
Specification-gaming-the-flip-side-of-AI-ingenuity](https://deepmind.com/blog/article/Specification-gaming-the-flip-side-of-AI-ingenuity)

- Reward can be difficult to define, e.g., how to define “good driving” using a number?

- Where does reward even come from?
- Human preferences? “Deep reinforcement learning from human preferences” by Christiano et al, 2017
- As a general rule, it's better to design performance measures according to what one actually wants to be achieved in the environment, rather than according to how one thinks the agent should behave. Safety -> car in parked in garage

RL Agent

- Agent: the decision-making process
- Environment: the “physical” components
- The agent–environment boundary represents the limit of the agent’s absolute control
- For convenience, people use the term *agent* loosely, e.g., we call robot *agent* even its physical body belongs to the environment

■ Policy


- ▶ deterministic: $\pi(S) = A$
- ▶ stochastic: $\pi(A|S) = P[A_t = A|S_t = S]$

■ Value function


- ▶ estimated total future reward starting from a state or a state-action pair
- ▶ help choose action (implicit policy)


■ Model (optional)


- ▶ agent's understanding of the environment dynamics (i.e., reward function and transition function)



- Rewards: -1 per time-step
- Actions: N, E, S, W
- States: Agent's location








- Model-based learning: building *models* of the environment, e.g., compute expectation by first estimate event probabilities and then multiply them with the corresponding function values.


$$E[f(x)] \approx \sum \hat{p}(x) f(x).$$

- Model-free learning: do not build *models* of the environment, e.g., compute expectation as the average of the sampled function values:


$$E[f(x)] \approx \frac{1}{n} \sum_{i=1}^n f(x_i).$$



- Model-based
 - ▶ Value function and/or policy
- Model-free
 - ▶ Value function and/or policy



- Value-based
 - ▶ Value function (no implicit policy)
- Policy-based
 - ▶ Policy (no value function)
- Actor-critic
 - ▶ Value function and Policy



- Learn optimal behavior via trial and error, based on environment's feedback



Brief History of RL

- Three threads joined in 80's to form the modern RL
 - ▶ Thread 1: animal learning (psychology)
 - ▶ Thread 2: optimal control (engineering)
 - ▶ Thread 3: artificial intelligence (computer science)

- “Dynamic Programming” by Bellman, 1957
- “Dynamic Programming and Stochastic Control” by Bertsekas, 1976
- “Neuro-Dynamic Programming” by Bertsekas and Tsitsiklis, 1996
 - ▶ Then “approximate dynamic programming”
 - ▶ Now “deep reinforcement learning”

MATH 7521 - ADP Stoch Optim & Control



MATH 7521 - ADP Stoch Optim & Control

(3) Basic concepts and mathematical foundations of neural networks, learning, nonlinear optimization and control. Exact and approximate optimization of the utility function. Bellman equation, approximate Bellman equation for solving multivariate optimization problems in real time. Partially observable variables, with random noise and tactical objectives varying in time. COREQUISITE(S): **NURS 7902**.

- World's first successful checkers-playing program by Arthur Samuel, 1959
- TD-GAMMON by Tesauro (among the world's best backgammon players), 1992
- Machine intelligence based on animal learning psychology by Harry Klopff, 1970's
- Temporal-difference learning, 1980's
- "Reinforcement Learning: An Introduction" by Sutton and Barto, 1998, (2nd ed. in 2020)

- Breakthroughs from DeepMind, OpenAI, etc.: Atari games (2015), AlphaGo (2016)...
- Mostly due to increased computing power + engineering practices, not much on the theory side
- Richard Bellman: “Once working in the area, it is very quickly realized that far more ability and sophistication is required to obtain a numerical solution than to establish the usual existence and uniqueness theorems”