

Generative Adversarial Network

Weizi Li

Department of Computer Science
University of Memphis



- Generative Adversarial Network
- Conditional GAN
- Example Code

Generative Adversarial Network

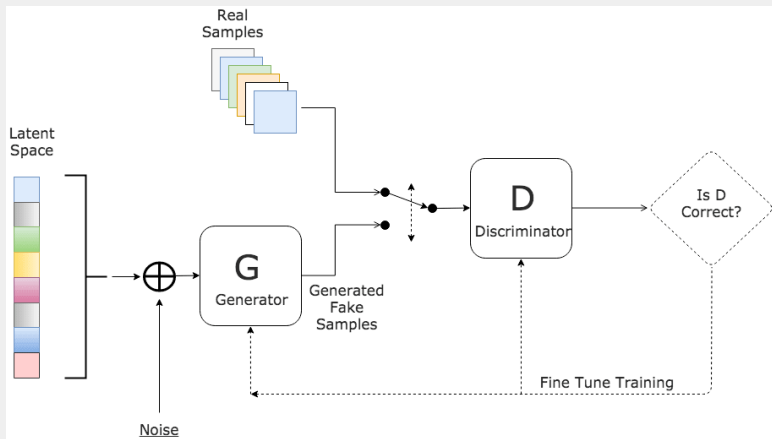
- Motivation: if deep neural network is so expressive, how about put two in competition?

- Goodfellow et al., Generative Adversarial Nets, 2014
- Generate synthetic data: images, sentences, ...
- “the most interesting idea in the last 10 years in Machine Learning” – Yann LeCun, Turing Award winner (2018)

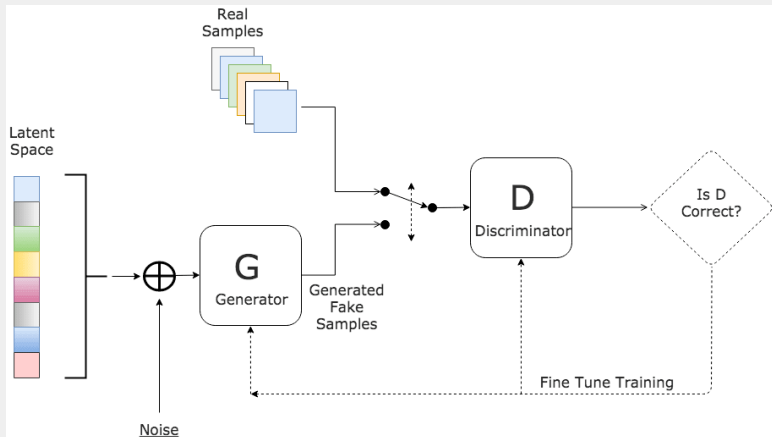
- Karras et al., A Style-Based Generator Architecture for Generative Adversarial Networks, 2019




- Generator (G): latent space \rightarrow synthetic data
- Each element in the latent space can represent a characteristic of the synthetic data



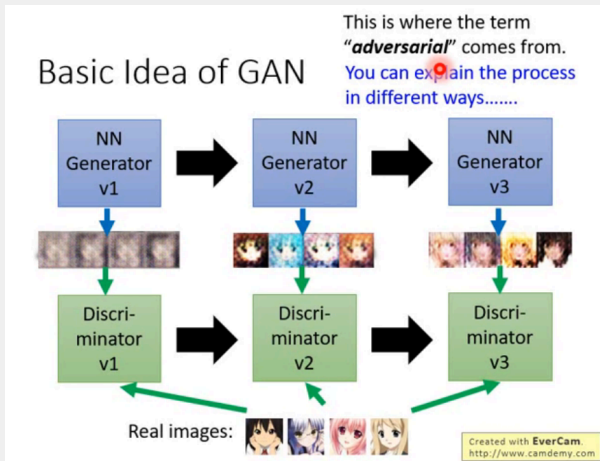
- Discriminator (D): synthetic data \rightarrow score (scalar)
- The score is used to judge whether the synthetic data is “realistic” or not



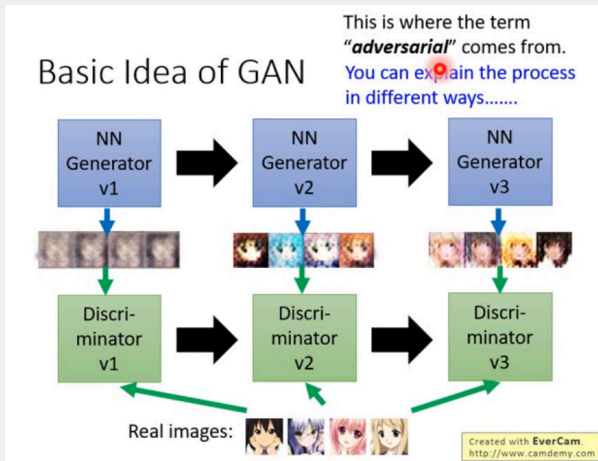
- To start: Generator v1 generates some synthetic data and Discriminator v1 tries to differentiate real and synthetic



- Generator v2's goal is to fool Discriminator v1 by improving synthetic data quality

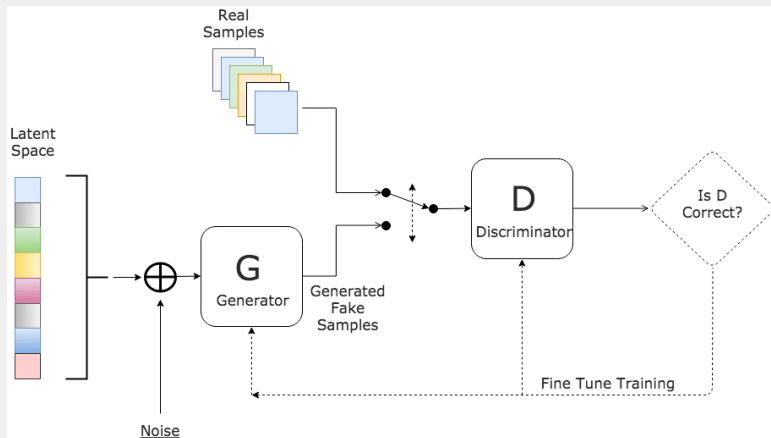


- Discriminator v2's goal is to defeat Generator v2 by identifying what is real and what is synthetic



- Step 1: initialize the parameters of generator and discriminator
- Step 2: fix generator, use it to generate synthetic data
- Step 3: sample real data from training set
- Step 4: **update** discriminator so that it can differentiate real and synthetic
- Step 5: fix discriminator, **update** generator till the current discriminator can no longer differentiate real and synthetic
- Step 6: go to Step 2

- During training, generator and discriminator are combined to form a large network with one hidden layer in the middle representing the generated data



Algorithm Initialize θ_d for D and θ_g for G

- In each training iteration:

Learning
D

- Sample m examples $\{x^1, x^2, \dots, x^m\}$ from database
- Sample m noise samples $\{z^1, z^2, \dots, z^m\}$ from a distribution
- Obtaining generated data $\{\tilde{x}^1, \tilde{x}^2, \dots, \tilde{x}^m\}$, $\tilde{x}^i = G(z^i)$
- Update discriminator parameters θ_d to maximize
 - $\tilde{V} = \frac{1}{m} \sum_{i=1}^m \log D(x^i) + \frac{1}{m} \sum_{i=1}^m \log (1 - D(\tilde{x}^i))$
 - $\theta_d \leftarrow \theta_d + \eta \nabla \tilde{V}(\theta_d)$

Learning
G

- Sample m noise samples $\{z^1, z^2, \dots, z^m\}$ from a distribution
- Update generator parameters θ_g to maximize
 - $\tilde{V} = \frac{1}{m} \sum_{i=1}^m \log (D(G(z^i)))$
 - $\theta_g \leftarrow \theta_g + \eta \nabla \tilde{V}(\theta_g)$

- [link](#)

Conditional GAN

- Conditional GAN represents a class of useful GAN variations: supervised, unsupervised

Text-to-Image

a dog is running



a bird is flying



- Traditional supervised approach

c^1 : a dog is running

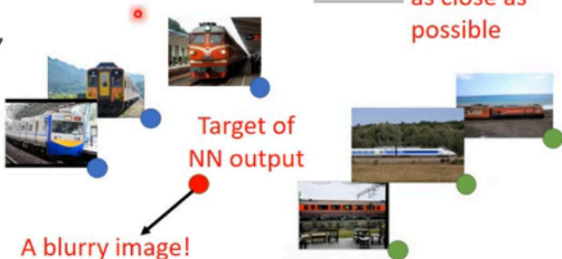
NN

Image




as close as possible



Text: "train"



Conditional GAN



True text-image pairs: (train, ) 1

(cat, ) 0 (train, ) 0


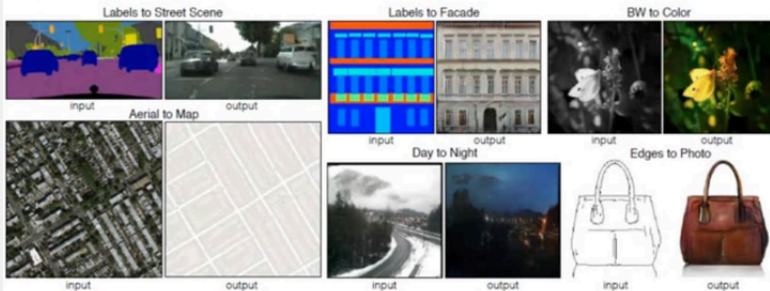



Image-to-image



<https://arxiv.org/pdf/1611.07004>

Image-to-image



- Traditional supervised approach



Testing:


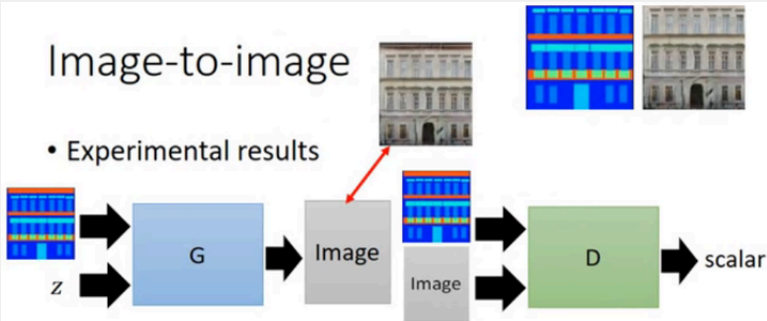


Image-to-image

- Experimental results

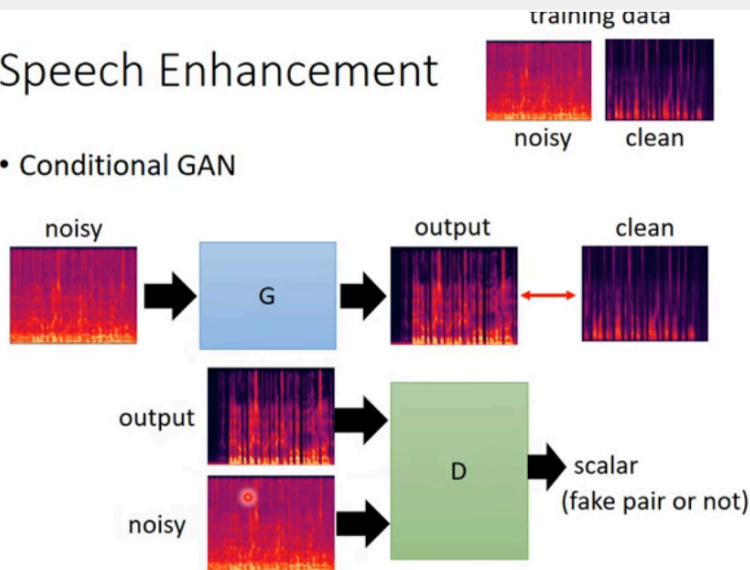


Testing:



Speech Enhancement

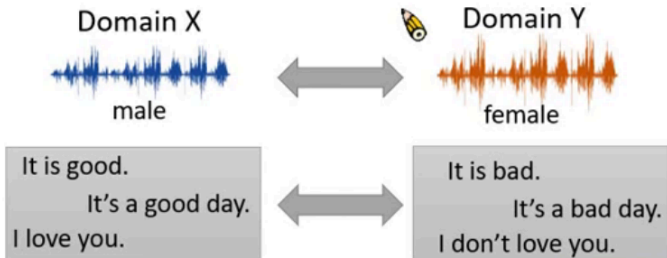
- Conditional GAN



Unsupervised Conditional Generation

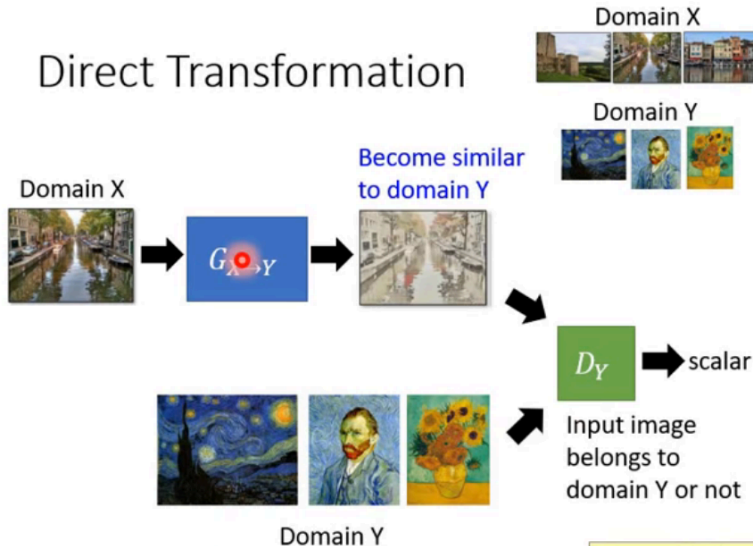


Transform an object from one domain to another
without paired data (e.g. style transfer)



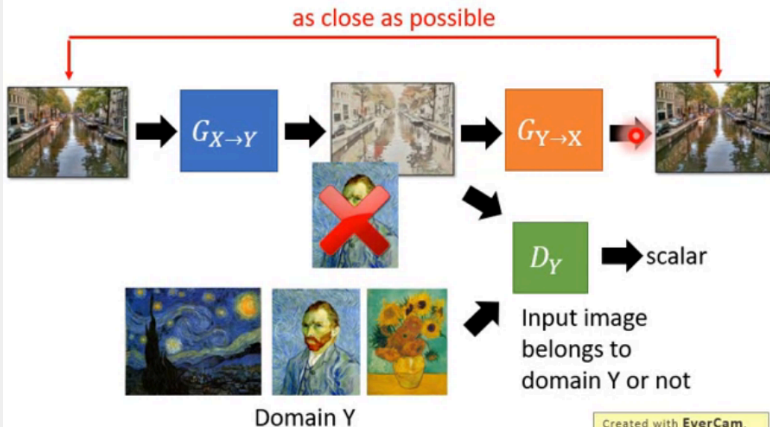
- Application: voice style transfer
- Past: need speech of the *same content* from two people
- Now: only need speech from two people

Direct Transformation



[Jun-Yan Zhu, et al., ICCV, 2017]

Direct Transformation



Example Code

- Conditional GAN: [link](#)